

Playing The Strengths of Synthetic Data

Carson Trego

September 2024

1 Introduction

In recent years, we have seen neural networks do a wealth of varied tasks from object detection, code generation, and sentiment analysis [3][1][2]. Neural networks, for all of the flexibility they can provide, often require significantly more training data compared to other machine learning methods [6]. Many large datasets are available for those seeking to train their own neural networks on established tasks, but if a large dataset cannot be used for a user's task, they will have to create a new dataset for said task [3][6]. Creating a high quality dataset is challenging: dataset creators will have to gather a large number of samples, annotate those samples mostly by hand, and those samples must have a fair distribution of characteristics, or the model could be biased. For example, one model gave the appearance of being able to classify huskies and wolfs, but because the backgrounds for wolfs consistently contained snow, the model would classify the subject as a wolf instead of a husky [2]. Gathering and annotating data for a neural network dataset is not only time and labor intensive, but ensuring fairness of the samples and preventing the wrong features from being used in prediction is a difficult task, so it is easy to put a large amount of work into a dataset that has a fatal flaw of some kind [6][2]. One proposed method to make datasets faster to create and more fair is using synthetic data to train a model [6][3][1][7][11]. Synthetic data is broad term to describe data that is simulated instead of being gathered directly from reality [1][11]. For example a classifier model may need to differentiate between two types of cups. Rather than obtaining a wealth of images of the cups in a variety of environments, a set of images could be rendered from the 3D models of the cups and then used to train the model. Ideally, a 3D rendered image should be so realistic that the model learns as effectively as it does with real images. In practice, there is a well-known issue called the "domain gap" that often occurs when synthetic data is used, in which the model performs well on the virtual training, validation, and testing data, but performs relatively worse when given real data [4][3]. The domain gap does not mean that the model does not perform well, just worse than the real data model [4]. Advances in rendering software could bring the domain gap closer, but there is still advantages to using synthetic data for a model, particularly when real data is either limited or completely unavailable otherwise [3][4][1]. Advantages to synthetic data methods are not just in cost

and effort, as that extra information, such as a 3D model of the object that is intended to be classified, can be used to do more with data that is available, as a 3D model has some information that is often not available from photos alone, and synthetic methods in general could provide more flexibility in classification, such as allowing for different classes that are slight modifications of the previous ones.

2 Theory

The method intended for this project begins with a 3D model of the objects that are intended to be classified. While classifying objects that vary significantly could be a topic that could be an aspect of this project, we will begin with assuming the simple case where each object is consistent in form, is fairly rigid, and opaque.

The first advantage in using synthetic data is that it eliminates a lot of redundancy in gathering data. Say a model needs to classify 8 different cups, which are manufactured to be nearly identical, so a cup that is in class 1 would look almost the exact same as a cup that is also in class one (As is true with disposable cups, which are often sold in large packs of identical cups). With the traditional method, several photos, at several angles, distances, lighting, and environments would have to be obtained, which would be a large effort to describe just one object. Using synthetic data and rendering system, a single model for each cup would be sufficient.

3D models contain information on how the object would look from all angles, so one 3D model can be used to simulate thousands of photos capturing different angles and distances. Assuming that a 3D render of an object is sufficiently realistic, the object could be observed in a variety of different chosen simulated situations with only one 3D model. The advantage with choosing the environments is that there is no background bias issue. Images of coconuts in tropical areas might be more common than images coconuts of in factories, but with synthetic data, a classifier could be trained in rare or nonsensical situations such as an image with a coconut in a steel mill. This ensures that, provided the rendering system has a wealth of virtual spaces, the model does not have a classification bias based on background features. The images in a synthetic data system can be easily balanced to represent each class equally and in varied environments.

The size of the object is another piece of information provided in 3D models that can be used. While normal photos often are not paired with specific information of the object, such as the angle the object is facing relative to some central point, the distance from the camera, or the actual size of the object, a 3D render has all of this information by design. If the model is then used with some sort of depth seeing method, such as a range finder, LIDAR, or stereo camera, the model could know how far away the object is, and combined with the known size of the class, “know” when a larger version of one class is being shown. This could also be achieved without depth-seeing hardware, provided the camera is

placed at a consistent distance from the objects, and is given information about how far the platform is. This aspect of synthetic data directly leads to 6 degree of freedom pose estimation, which is very useful for robotics applications. This will be explored in greater depth in a later part of the document.

Black box machine learning algorithms do not allow the operators to predict how the model will respond to certain conditions, and the method of finding out how a model will respond to specific environments is by testing the model directly. This adds another step to the model making process, where a model must first be completed in entirety and then tested under a variety of conditions. Using synthetic data, this aspect of creating models is not fully removed, but the method of producing new data to test the model on can be used to predict how the model will fail, and potentially adjust for it. For example, a model may struggle to classify two objects that look identical from certain angles. Using synthetic data, the model could automatically test itself on all of the classes to check for weak points, and “know” before hand that it will struggle with the class at certain angles, and when the model is presented with a said object at said angle, it can instruct itself to obtain more information on the object, rather than using the current image to make a definitive classification.

Theoretically, synthetic data in the form of 3D models allows for single models of objects to provide images from all angles, a wealth of environments, information in 6 degree of freedom pose estimation, the size of the object, and self diagnostics for when the model has shortcomings. This all assumes a strong 6 degree of freedom pose estimation and classification model to begin with, alongside a high quality rendering system with a wealth of environmental conditions prepared in advance. With these assumptions, there should be no domain gap, yet there is, so while the information needed to make these conclusions is present in the system, and while a human operator given a 3d model could theoretically perform this task in entirety, how the model responds to the actual synthetic data, and the quality of the synthetic data, is to be seen.

3 Defining The Challenge

The goal of this challenge is to exploit the advantages of synthetic data, in which all objects are placed on a platform which the model can make observations of through a camera. The model should “know” how far away the platform is, and be able to rotate the platform and gather additional pictures if needed. To quantify the model’s performance, the model will have to correctly classify the object and its 6 degree of freedom parameters in its final conclusion, as well as do so with as few photos and mechanical actions as possible.

To demonstrate this we will give two examples in a model with 3 classes: Flathead screw, Phillips head screw, and plastic banana. In the first case, the model identifies that the object is a plastic banana and recognizes that the probability of misclassifying a banana from this angle is rare, so it automatically logs the class and 6 degree of freedom parameters. In the second example, the model sees a screw with the head facing the opposite direction. Having

prior tests show that a screw of either type at this angle is likely to result in a misclassification, the model calculates what angles of the objects are most likely to provide a definitive result between the two, by finding angles where the correct classification rate is high for both objects. Using this information in the form of a spherical space, the model runs a second calculation to find the area that would function as a tie breaker that is the closest to the current angle (thus reducing the amount of mechanical effort and time needed), and snaps an additional photo of the object. This next photo (and potentially prior information) is used to make a final classification of the object and its 6 degree of freedom parameters.

4 Other Efforts

Many studies and model creators have used synthetic data as a way to augment their dataset or create a dataset from scratch [3][4][6][11][7]. 3D models have been used as synthetic data for 6 degree of freedom pose estimation, as in the case of the LINEMOD template matching method [9][10]. There is also 6D models such as EfficientPose and YOLO6D, which both perform 6 degree of freedom pose estimation using 3D models, but make use of a convolutional neural network. [13][5]. In a prior project, I had used a method similar to the copy paste method, but rather than using segmented real images, I had taken 3D models and segmented them onto a different background, this method was later used in plain classification and with other models [8][12][14]. While the object detection models were able to perform quantifiable well on the virtual data, the real life tests were limited, but showed fair results [12][14]. Much of the real world testing has been fairly qualitative, however the real time capability of the model allow for a better intuition of its performance. For this project, a classifier and pose estimator could be used independently to improve accuracy, and potentially, the amount of operations required for processing an object.

These model articles tend to not discuss a method for scaling its use, and the goal of this project is to bring these models closer to being part of a system that could be efficiently used in a variety of devices. One aspect to make the models more usable will be the efficient tie breaking method, in which the model prepares angles and objects that tend to produce confusion, and how to most efficiently determine which class is being shown. The other aspect of this model is physical, as the platform that is used for the objects and the method of scanning new objects and automatically adjusting the model to accommodate them will be part of the model considerations.

5 Method

Several objects will be prepared as 3D model and physical object pairs obtained by either scanning a physical object or 3D printing a 3D model. The 3D models will be used with a synthetic data generator to produce training data and tem-

plates, which will then be used in the 6 degree of freedom pose estimator and the classifier. Once the models are prepared, the synthetic data generator will be used to access potential weak points of the model by attempting to classify the object at a multitude of angles and conditions. These results will be mapped to a spherical plane and saved for future reference. When performing the physical test, these saved tests and the pose estimator will be used to estimate the risk of misclassification caused by having an angle that does not offer enough information to adequately classify the object, and the classification test spheres will be used to find positions that would provide more information on the object, while comparing the results to other objects that are at risk, with a set of points on the spheres chosen to maximize the chance of a decisive classification of the object. In addition to searching for regions that would provide the most information, the model will be attempting to find the shortest path and least number of photos required to come to a decisive conclusion. This test will be ran several times on multiple different objects, and the amount of operations, camera movement, and correct classifications will all be incorporated into how well the overall system performs. This should be more accurate and faster than taking a large amount of photos and processing all of them, or taking photos of the object uniformly at random, but this remains unknown. The goal of this challenge is to see if applying this method provides more accuracy and speed than simply using an object detector on one photo alone, grid searching the object processing a large number of photos, or randomly taking photos of the object and processing those.

References

- [1] What is synthetic data? URL: <https://www.ibm.com/topics/synthetic-data>.
- [2] Husky or wolf? using a black box learning model to avoid adoption errors, 2017. URL: <https://innovation.uci.edu/2017/08/husky-or-wolf-using-a-black-box-learning-model-to-avoid-adoption-errors/>.
- [3] Bridging the domain gap for neural models, 2019. URL: <https://machinelearning.apple.com/research/bridging-the-domain-gap-for-neural-models>.
- [4] Xiangyu Bai, Yedi Luo, Le Jiang, Aniket Gupta, Pushyami Kaveti, Hanumant Singh, and Sarah Ostadabbas. Bridging the domain gap between synthetic and real-world data for autonomous driving, 2023. URL: <https://arxiv.org/abs/2306.02631>, <https://arxiv.org/abs/2306.02631> arXiv:2306.02631.
- [5] Yannick Bukschat and Marcus Vetter. Efficientpose: An efficient, accurate and scalable end-to-end 6d multi object pose estimation approach, 2020. URL: <https://arxiv.org/abs/2011.04307>, <https://arxiv.org/abs/2011.04307> arXiv:2011.04307.

- [6] Terrance DeVries and Graham W. Taylor. Dataset augmentation in feature space, 2017. URL: <https://arxiv.org/abs/1702.05538>, <https://arxiv.org/abs/1702.05538> arXiv:1702.05538.
- [7] Killeen B. D. Hu Y. Grupp R. B. Taylor R. H. Armand M. Unberath M. Gao, C. Synthetic data accelerates the development of generalizable learning-based algorithms for x-ray image analysis. *Nature machine intelligence*, 2023. URL: <https://doi.org/10.1038/s42256-023-00629-1>.
- [8] Golnaz Ghiasi, Yin Cui, Aravind Srinivas, Rui Qian, Tsung-Yi Lin, Ekin D. Cubuk, Quoc V. Le, and Barret Zoph. Simple copy-paste is a strong data augmentation method for instance segmentation, 2021. URL: <https://arxiv.org/abs/2012.07177>, <https://arxiv.org/abs/2012.07177> arXiv:2012.07177.
- [9] Stefan Hinterstoisser, Vincent Lepetit, Slobodan Ilic, Stefan Holzer, Gary Bradski, Kurt Konolige, and Nassir Navab. Linemod dataset, 2013.
- [10] Stefan Hinterstoisser, Vincent Lepetit, Slobodan Ilic, Stefan Holzer, Gary Bradski, Kurt Konolige, and Nassir Navab. Model based training, detection and pose estimation of texture-less 3d objects in heavily cluttered scenes. In Kyoung Mu Lee, Yasuyuki Matsushita, James M. Rehg, and Zhanyi Hu, editors, *Computer Vision – ACCV 2012*, pages 548–562, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg.
- [11] German Ros, Laura Sellart, Joanna Materzynska, David Vazquez, and Antonio M. Lopez. The synthia dataset: A large collection of synthetic images for semantic segmentation of urban scenes. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 3234–3243, 2016. <https://doi.org/10.1109/CVPR.2016.352> doi:10.1109/CVPR.2016.352.
- [12] Ben Laube Carson Trego Maxwell Kline Tyson Shields. Yosco, you only scan once. using a single 3d scan to automatically generate a labeled object detection training dataset, 2024.
- [13] Bugra Tekin, Sudipta N. Sinha, and Pascal Fua. Real-Time Seamless Single Shot 6D Object Pose Prediction. In *CVPR*, 2018.
- [14] Carson Trego. Yosco cls, 2024. temp, currently a program on its own with no written document, as this is part of that.